

Luke Anglin

AI Engineer | Virginia Beach, VA, USA | +17574781566 | lukeanglin@gmail.com | luketyleranglin.github.io



Professional Summary

I am an AI/ML Engineer with 4 to 5 years of experience deploying large-scale cloud-based production software, specializing in high-performance computing pipelines. I have experience managing massive inference protocols for image and video generation, using AWS Sagemaker, TensorRT, and Triton Inference Server. I manage the full-stack software development lifecycle, and my expertise includes programming languages like Python, TypeScript, Go, Rust, and C++, as well as a bit of Java and SQL.

I have a strong understanding of data modeling, and have led the migration of large banking clients to modern cloud architectures in Azure and Snowflake, enabling more efficient querying on large datasets and reducing reliance on multi-million dollar on-premise infrastructure.

I have extensive experience with data pipeline development, including building, maintaining, and optimizing ETL/ELT pipelines. I have a deep understanding of parallelism and asynchronous workflows, shown by my experience multithreading Kafka operations to enhance system efficiency at Tesla. I have leveraged tools like Docker, Jenkins, Terraform, and Prometheus for efficient CI/CD and strong monitoring/observability. I am also familiar with network engineering, and I have a certification as a CCNA, Cisco Certified Networking Associate.

I have worked with LLMs and led the development of agentic systems and foundational models in order to find savings through automation. Projects range from structured output extraction on long electronic medical records to agentic React-based web UIs for company executives, both using SOTA LLM models and using small, fine-tuned ones or ones developed with tiny pretraining runs. I understand quantization and know how to make inference and training efficient, ensuring strong GPU utilization and low latency or high throughput, depending on the context. I have additionally used RLHF and evals to improve agents and models.

Outside of my work, I spend large amounts of times testing robotic systems. I enjoy learning about ROS2, NVIDIA Isaac, and how IoT and edge systems interact with AI and software. I have a strong understanding of multimodal LLMs.

Experience

Scale AI

New York, NY, USA

Forward Deployed AI Engineer - Frontier Agents

May '26 - Present

- Built frontier agent technology in messy enterprise environments
- Built out continual learning environments for AI agents
- Provisioned cross-tenant capabilities and managed conversations with clients
- Leveraged RLHF and evals to improve agents and agent scaffolds

Adobe

San Jose, CA, USA

Forward Deployed AI Engineer

Nov '25 - Apr '26

- Worked with clients to implement agentic, human-in-the-loop systems for AI-centered marketing workflows
- Designed high-performance, high-volume auto scaling workflows for large-scale image generation, leveraging cloud GPUs

Elevance Health

AI Engineer

Norfolk, VA

Mar '24 - Nov '25

- Engineered a computer vision tool for EMR extraction, automating vendor work. ~\$2 million saving
- Leveraged PyTorch, CUDA, and NCCL to design an reinforcement learning (RL) model for the Medicare bid process
- Led the development of a MCP-enabled RAG application for our M&A team, reducing ~200 hours per project
- Designed and implemented a voice-responsive AI agent using OpenAI Whisper

EY

AI Technology Consultant

Charlotte, NC

Sep '23 - Mar '24

- Architected Langgraph-based applications to automate client's vendor work. ~\$4 million in savings
- Migrated banking clients from on-prem systems to Azure and Snowflake
- Wrote high-performance, modern C++ and Rust code to optimize legacy financial systems.

Tesla

Data Science Intern

Fremont, CA

May '22 - Aug '22

- Multithreaded Kafka operations to enhance system efficiency through concurrency.
- Utilized a suite of tools including Jenkins, Splunk, Postman, and Grafana to analyze logs and debug complex issues
- Worked on large, high-performance Golang repositories, leveraging CI/CD and source control to drive efficient mobile routing.

University of Virginia

Deep Learning Research Assistant

Charlottesville, VA

Sep '21 - Sep '22

- Employed inference and learning techniques across edge devices (mobile phones)
- Worked hands-on with GPUs; built federated data pipelines to optimize training

Elevance Health

Artificial Intelligence Team Intern

Palo Alto, CA

May '20 - Aug '20

- Designed digital twin AI systems for healthcare members

Education

University of Virginia

May '23

Bachelor of Science | Computer Science | GPA: 3.9

Received full academic scholarship. Graduated in three years. Participated in UVA MARS (Mechatronics and Robotics) consistently, working on our lunar project.

Skills

HW/SW Codesign · Kubernetes (K8s) · High Performance Computing (HPC) · KV Cache · Open Source Contributor · MoE · vLLM · JAX · Android · IOS · React · Ray · Pretraining/Fine-tuning/Post-training · Diffusion · Multimodal AI · VLAs · World Models · Typescript/Javascript/Node · AWS/Azure/GCP · German · Statistics · Unreal Engine